


Systems and Devices 2 (SYS2) Module
Threads/Tasks Synchronization
Autumn Term 2021

Concurrency:

- Concurrency in OS is allowing more than one process to make progress at a time
- The basic requirement for support of concurrent processes is the ability to enforce mutual exclusion
- Example
 - Processes P_0 and P_1 are creating child processes using the `fork()` system call
 - Race condition on kernel variable `next_available_pid` which represents the next available process identifier, `pid`
 - Unless there is a mechanism to prevent P_0 and P_1 from accessing the variable `next_available_pid` the same `pid` could be assigned to two different processes

**Race condition:**

- A race condition occurs when multiple processes or threads read and write data items so that the final result depends on the order of execution of instructions in the multiple processes

Critical Section Problem:

- Consider system of n processes p_0, p_1, \dots, p_{n-1}
- Each process has critical section segment of code
- Process may be changing common variables, updating table, writing file
- When one process in critical section, no other may be in its critical section, no other may be in the critical section
- Critical section problem is to design protocol to solve this
- Each process must ask permission to enter critical section in entry section, may follow critical section with exit section, then remainder section

General structure of process P_i .

```

do {
    entry section
    critical section
    exit section
    remainder section
} while (true);
  
```

- Critical section problem solution requirements
 - Mutual exclusion
 - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
 - Progress
 - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the process that will enter the critical section next cannot be postponed indefinitely
 - Bounded waiting
 - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made the request to enter its critical section and before that request is granted
 - Assume that each process executes at a Nonzero speed
 - No assumption concerning relative speed of the n processes

Requirements for Mutual Exclusion:

- Mutual exclusion must be enforcedfi Only one process at a time is allowed into its critical section, among all processes that have critical section among all processes that have critical sections for the same resource or shared object
- A process that halts in its noncritical section must do so without interfering with other processes
- It must not be possible for a process requiring access to a critical section to be delayed indefinitely, no deadlock or starvation
- When no process is in a critical section any process that requests entry to its critical section must be permitted to enter without delay
- No assumptions are made about relative process speeds or number of processors
- A process remains inside its critical section for a finite time onlyfi

Implementation of Mutual Exclusion:

- One approach is to leave the responsibility with the processes that wish to execute concurrentlyfi Processes, whether they are system programs or application programs, would be required to coordinate with one another to enforce mutual exclusion, we can refer to these as software approachesfi Although this approach is prone to high processing overhead and bugs, it is nevertheless useful to examine such approaches to gain a better understanding of the complexity of concurrent processingfi
- A second approach involves the use of special-purpose machine instructionsfi These have the advantage of reducing overhead but nevertheless will be shown to be unattractive as a general-purpose solution
- A third approach is to provide some level of support within the OS as a programming languagefi

Mutual exclusion – hardware support:

- Atomic operation

- A function or action implemented as a sequence of one or more instructions that appears to be indivisible, that is, no other process can see an intermediate state or interrupt the operation. The sequence of instructions is guaranteed to execute as a group, not execute at all, having no visible effect on system state. Atomicity guarantees isolation from concurrent processes.
 - Interrupt disabling
 - In a uniprocessor system, concurrent processes cannot have overlapped execution, they can only be interleaved. Furthermore, a process will continue to run until it invokes an OS service or until it is interrupted. This capability can be provided in the form of primitives defined by the OS kernel for disabling and enabling interrupts.

Special Machine Instructions:

- In a multiprocessor configuration, several processors share access to a common main memory. At the hardware level, access to a memory location excludes any other access to that same location.
- With this as a foundation, processor designers have proposed several machine instructions that carry out two actions atomically, such as reading and writing or reading and testing, of a single memory location with one instruction fetch cycle. During execution of the instruction, access to the memory location is blocked for any other instruction referencing that location.

Special Machine Instructions:

Compare & Swap Instruction

```
int compare_and_swap (int *word, int testval, int newval)
{
    int oldval;
    oldval = *word;
    if (oldval == testval) *word = newval;
    return oldval;
}
```

Exchange Instruction

```
void exchange (int *register, int *memory)
{
    int temp;
    temp = *memory;
    *memory = *register;
    *register = temp;
}
```

Advantages of Machine-instruction approach:

- It is applicable to any member of processes on either a single processor or multiple processors sharing main memory
- It is simple and therefore easy to verify
- It can be used to support multiple critical sections, each critical section can be defined by its own variable

Disadvantages of Machine-Instruction Approach:

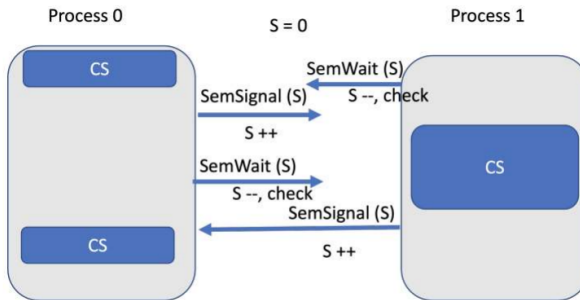
- Busy waiting is employed
 - While a process is waiting for access to a critical section, it continues to consume processor time
- Starvation is possible
 - When a process leaves a critical section and more than one process is waiting, process is arbitrary. Some process could indefinitely be denied access
- Deadlock is possible
 - Consider the following scenario on a single-processor system. Process P1 executes the special instruction (e.g., compare and swap, exchange) and enters its critical section. P1 is then interrupted to give the processor to P2, which has higher priority. If P2 now attempts to use the same resource as P1, it will be denied access because of the mutual exclusion mechanism. Thus, it will go into a busy waiting loop. However, P1 will never be dispatched because it is of lower priority than another ready process, P2.

OS and Programming language supported concurrency:

- The idea that two or more processes can cooperate by means of simple signals, such that a process can be forced to stop at a specified place until it has received a specific signal. Any complex coordination requirement can be satisfied by the appropriate structure of signals. Different data structures that can be used are, Semaphore
- Binary Semaphore, Mutex, Conditional Variables, Event Flags, Mailbox messages, spinlocks

Semaphore:

- An integer value used for signalling among processes
- Only three operations may be performed on a semaphore, all of which are atomic: initialize, decrement, and increment
- The decrement operation may result in the blocking of a process, and the increment operation may result in the unblocking of a process. Also known as a counting semaphore or a general semaphore
- Example



```
struct semaphore {
int count;
queueType queue;
};

void semWait(semaphore s){
s.count--;
if (s.count < 0) {
/* place this process in s.queue */;
/* block this process */;
}}

void semSignal(semaphore s){
s.count++;
if (s.count <= 0) {
/* remove a process P from s.queue */;
/* place process P on ready list */;
}}
```

Binary semaphore:

- A binary semaphore may be initialized to 0 or 1
- The semWaitB operation checks the semaphore value. If the value is zero, then the process executing the semWaitB is blocked. If the value is one, then the value is changed to zero and the process continues execution.
- The semSignalB operation checks to see if any processes are blocked on this semaphore. If this is the case, then a process blocked by a semWaitB operation is unblocked. If no processes are blocked, then the value of the semaphore is set to one.

```

struct binary_semaphore {
    enum {zero, one} value;
    queueType queue;
};
void semWaitB(binary_semaphore s){
    if (s.value == one)
        s.value = zero;
    else {
        /* place this process in s.queue */;
        /* block this process */;
    }
}
void semSignalB(semaphore s){
    if (s.queue is empty())
        s.value = one;
    else { /* remove a process P from s.queue */;
        /* place process P on ready list */;
    }
}

```

Mutex:

- Similar to a binary semaphore
- A key difference between the two is that the process that locks the mutex, sets the value to 0, must be the one to unlock it, sets the value to one

Monitor and conditional variables:

- A monitor is a software module consisting of one or more procedures, an initialization sequence, and local data. The chief characteristics of a monitor are the following
 - The local data variables are accessible only by the monitor's procedures and not by any external procedure
 - A process enters the monitor by invoking one of its procedures
 - Only one process may be executing in the monitor at a time, any other processes that have invoked one of its procedures
 - Only one process may be executing in the monitor at a time, any other processes that have invoked the monitor are blocked, waiting for the monitor to become available

Condition variables:

- Condition variables are a special data type in monitors, which are operated on by two functions
 - Cwait
 - Suspend the execution of the calling process on condition c. The monitor is now available for use by another process
 - Csignal

- Resume execution of some process blocked after a cwait on the same condition. If there are several such processes, choose one of them; if there is no such process, do nothing

